

RISC-V Instruction Encoding - 1



Z. Jerry Shi

Department of Computer Science and Engineering
University of Connecticut

CSE3666: Introduction to Computer Architecture

Outline

- RISC-V Instruction encoding
 - Instruction format
 - Encoding of instructions like ADD, ADDI, and Load/store
 - Decoding

Design Principle 4: Good design demands good compromises

Keep formats as similar as possible

Reading: Section 2.5 and the beginning of Section 2.10.

References: Reference card in the book.

Representing instructions with bits


- We use bits to represent numbers, characters, etc.
- We also use bits to represent instructions

Design questions:

- How many bits should we use to encode instructions?
- Are we using the same number of bits to encode all instructions?
 - Do all instructions have the same length?

RISC-V instruction words

- RISC-V base ISA are encoded as 32-bit **instruction words**
 - Encoded instructions are also called machine (language) code
- Both instructions and data are stored in memory
- **Program Counter (PC)** points to the current instruction
 - Incremented by 4 in normal flow for **sequential execution**



A red arrow labeled "PC" points to the row in the table where the memory address is $x + 4$ and the instruction is "Instr 11".

| Memory Address | Instructions |
|----------------|--------------|
| $x + 12$ | Instr 13 |
| $x + 8$ | Instr 12 |
| $x + 4$ | Instr 11 |
| x | Instr 10 |
| $x - 4$ | ... |

C extension allows compressed instructions of 16 bits, but it is not a stand-alone ISA. **Machine language** uses binary representation of instructions. Instructions in machine language are called **machine code**.

Discussion

- What information do you want to keep in the instruction word?

ADD rd, rs1, rs2

ADDI rd, rs1, imm

LW rd, offset(rs1)

SW rs2, offset(rs1)

instruction

32 bits

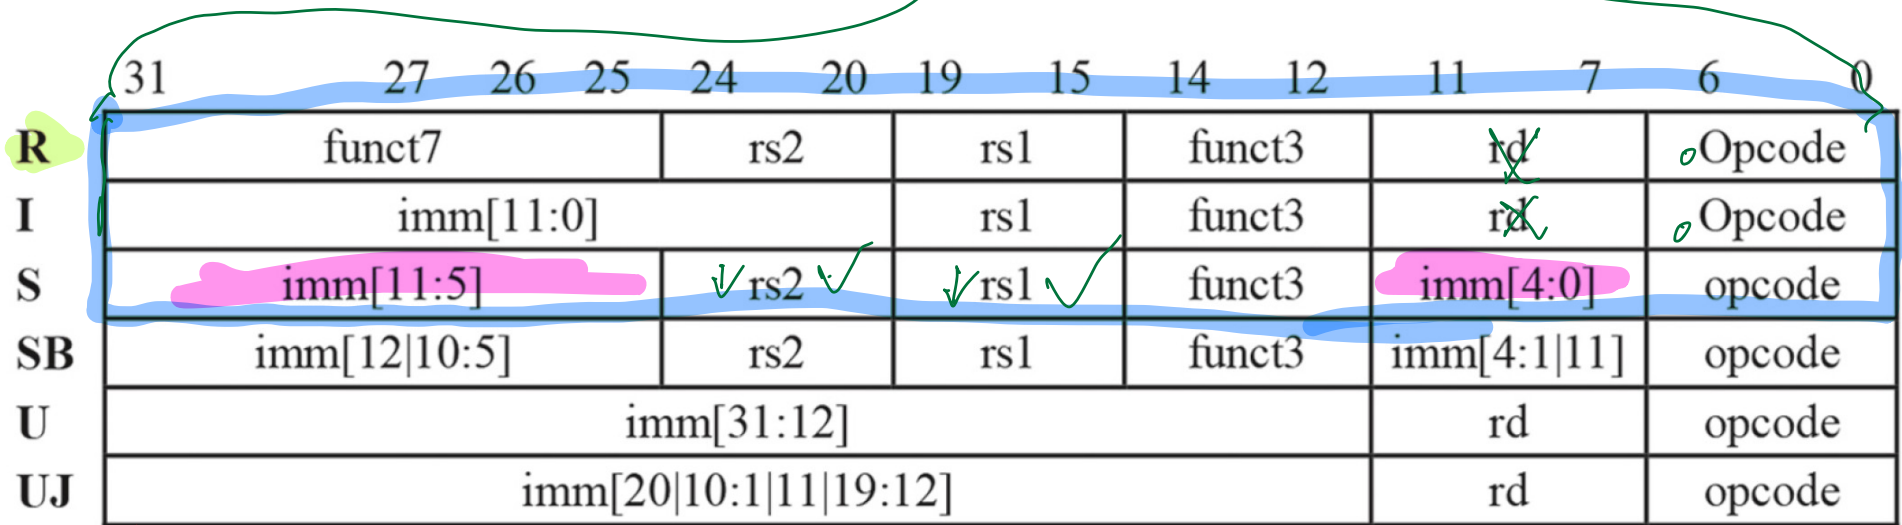
Instruction format

- The layout of bits in instruction words is **instruction format**
 - How do we use 32-bit bits to specify operation code (opcode), registers, immediate, offset, etc? How many bits for each?
- RISC-V has six instruction formats (while MIPS has 3)
 - R, I, S, SB, U, and UJ *Inst $\xrightarrow{?}$ machine code*
- Instructions we have learned so far
 - Using registers only
 - Having an immediate as the second operand
 - Load and store
 - LUI and Branch, to be discussed next week

Binary compatibility allows compiled programs to work on different computers

RISC-V Core Instruction Format

32-bit machine code



| denotes concatenation

The (green) card in the textbook

RISC-V Reference Data Card ("Green Card")

slt R Set Less Than $R[rd] = (R[rs1] < R[rs2]) ? 1 : 0$
 slti I Set Less Than Immediate $R[rd] = (R[rs1] < imm) ? 1 : 0$
 sltiu I Set < Immediate Unsigned $R[rd] = (R[rs1] < imm) ? 1 : 0$
 sltu R Set Less Than Unsigned $R[rd] = (R[rs1] < R[rs2]) ? 1 : 0$
 sra, sraw R Shift Right Arithmetic (Word) $R[rd] = R[rs1] >> R[rs2]$
 srai, sraw I Shift Right Arithmetic (Word) $R[rd] = R[rs1] >> imm$
 srl, srlw R Shift Right (Word) $R[rd] = R[rs1] >> R[rs2]$
 srli, srlw I Shift Right Immediate (Word) $R[rd] = R[rs1] >> imm$
 sub, subw R SUBtract (Word) $R[rd] = R[rs1] - R[rs2]$
 sw S Store Word $M[R[rs1] + imm(31:0)] = R[rs2](31:0)$
 xor R XOR $R[rd] = R[rs1] \oplus R[rs2]$
 xori I XOR Immediate $R[rd] = R[rs1] \oplus imm$

Notes: 1) The Word version only operates on the rightmost 32 bits of a 64-bit registers
 2) Operation assumes unsigned integers (instead of 2's complement)
 3) The least significant bit of the branch address in jalr is set to 0
 4) (signed) Load instructions extend the sign bit of data to fill the 64-bit register
 5) Replicates the sign bit to fill in the leftmost bits of the result during right shift
 6) Multiply with one operand signed and one unsigned
 7) The Single version does a single-precision operation using the rightmost 32 bits of a 64-bit F register
 8) Classify writes a 10-bit mask to show which properties are true (e.g., -inf, -0, +0, +inf, denorm, ...)
 9) Atomic memory operation; nothing else can interpose itself between the read and the write of the memory location
 The immediate field is sign-extended in RISC-V

CORE INSTRUCTION FORMATS

| | 31 | 27 | 26 | 25 | 24 | 20 | 19 | 15 | 14 | 12 | 11 | 7 | 6 | 0 |
|----|-----------------------|----|----|----|----|-----|----|--------|----|--------|----|-------------|--------|--------|
| R | funct7 | | | | | rs2 | | rs1 | | funct3 | | rd | | Opcode |
| I | imm[11:0] | | | | | rs1 | | funct3 | | rd | | Opcode | | |
| S | imm[11:5] | | | | | rs2 | | rs1 | | funct3 | | imm[4:0] | | opcode |
| SB | imm[12 10:5] | | | | | rs2 | | rs1 | | funct3 | | imm[4:1 11] | | opcode |
| U | imm[31:12] | | | | | | | | | | rd | | opcode | |
| UJ | imm[20 10:1 11 19:12] | | | | | | | | | | rd | | opcode | |

© 2018 by Elsevier, Inc. All rights reserved. From Patterson and Hennessy, Computer Organization and Design: The Hardware/Software Interface: RISC-V Edition

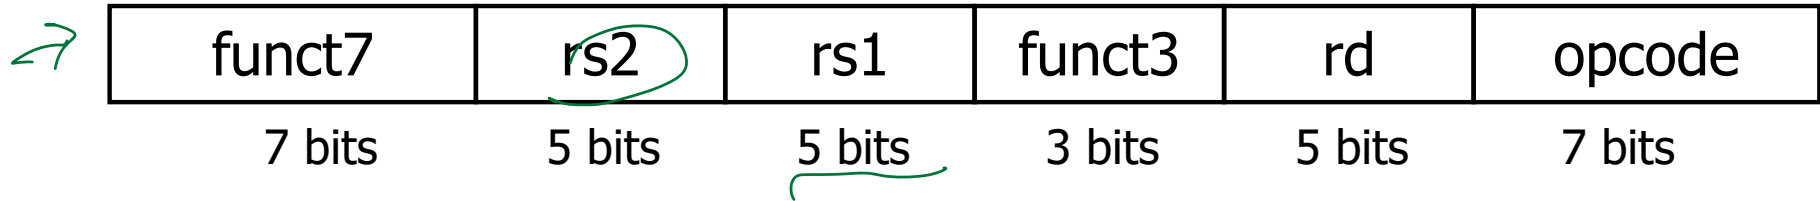
imm is the 32-bit immediate processor uses for computation.
 It is not the immediate written in instruction.

RISC-V R-type Instructions

| | | | | | |
|--------|--------|--------|--------|--------|--------|
| funct7 | rs2 | rs1 | funct3 | rd | opcode |
| 7 bits | 5 bits | 5 bits | 3 bits | 5 bits | 7 bits |

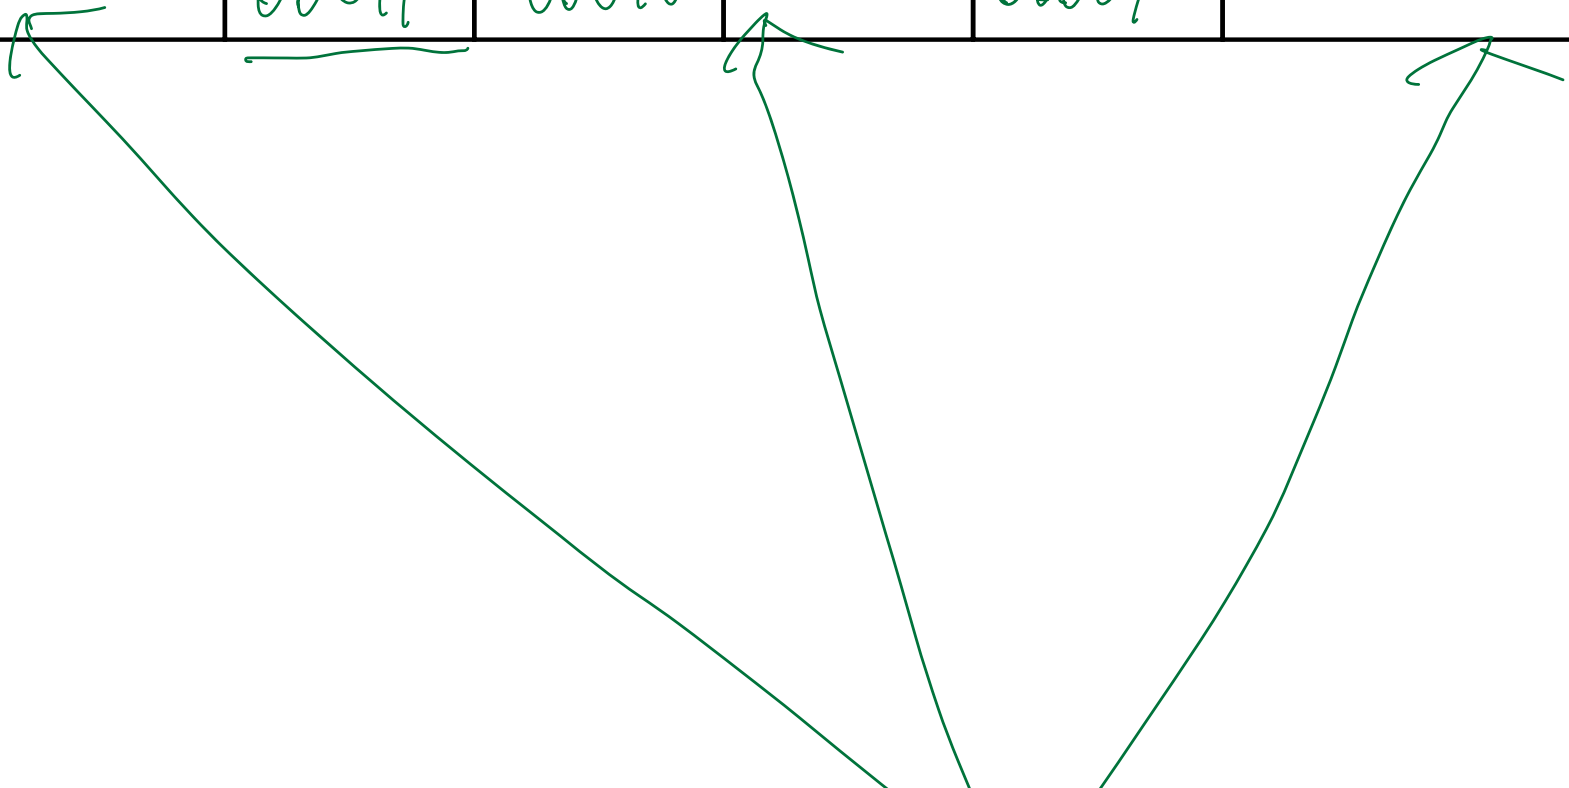
- For instructions that have 3 registers as operands
- Fields in R-type
 - opcode: 7-bit operation code
 - rd: destination register number
 - rs1: first source register number
 - rs2: second source register number
 - funct3: additional function code
 - funct7: even more function code

R-format Example: ADD



add x1, x2, x3

add Rd, Rs1, Rs2
x1, x2, x3



Opcode and funct codes

- From the green card, which also has hexadecimal representation

| | type | opcode | funct3 | funct7 |
|-------|------|----------------|------------|----------------|
| → add | R | <u>0110011</u> | <u>000</u> | <u>0000000</u> |
| sub | R | 0110011 | 000 | 0100000 |
| sll | R | 0110011 | 001 | 0000000 |
| slt | R | 0110011 | 010 | 0000000 |
| sltu | R | 0110011 | 011 | 0000000 |
| xor | R | 0110011 | 100 | 0000000 |
| srl | R | 0110011 | 101 | 0000000 |
| sra | R | 0110011 | 101 | 0100000 |
| or | R | 0110011 | 110 | 0000000 |
| and | R | 0110011 | 111 | 0000000 |

The right most two bits in opcode are always 11
for 32-bit instructions

R-type Example: ADD

| funct7 | rs2 | rs1 | funct3 | rd | opcode |
|--------|--------|--------|--------|--------|--------|
| 7 bits | 5 bits | 5 bits | 3 bits | 5 bits | 7 bits |

add x1, x2, x3

| | | | | | |
|---|---|---|---|---|------|
| 0 | 3 | 2 | 0 | 1 | 0x33 |
|---|---|---|---|---|------|

| | | | | | | | |
|----------|-------|-------|-----|-------|----------|---|---|
| 000 0000 | 00011 | 00010 | 000 | 00001 | 011 0011 | | |
| 0x0 | 0 | 3 | 1 | 0 | 0 | B | 3 |

0x 003100B3 ✓

1111 = F
0000 = 0

What if we change ADD to SUB? How many bits are changed?

Question (from textbook)

- What RISC-V instruction does this represent?

The table lists the bits in each field.

| funct7 | rs2 | rs1 | funct3 | rd | opcode |
|----------|-------|-------|--------|-------|----------|
| 010 0000 | 01001 | 01010 | 000 | 01011 | 011 0011 |

- A. sub x9, x10, x11
- B. add x11, x9, x10
- C. sub x11, x10, x9
- D. sub x11, x9, x10

| | Funct7 |
|-----|-------------|
| ADD | 0b 000 0000 |
| SUB | 0b 010 0000 |

RISC-V R-type Instructions



- We can use the format to encode any instructions like

`instr_name` `rd, rs1, rs2`

- How about instructions like `addi`, `slli`?

`addi` `rd, rs1, imm`

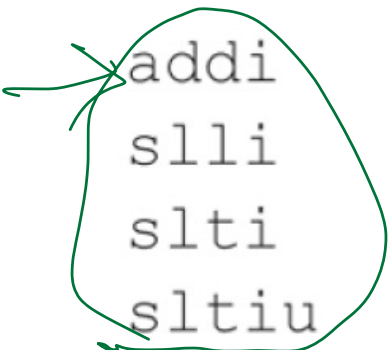
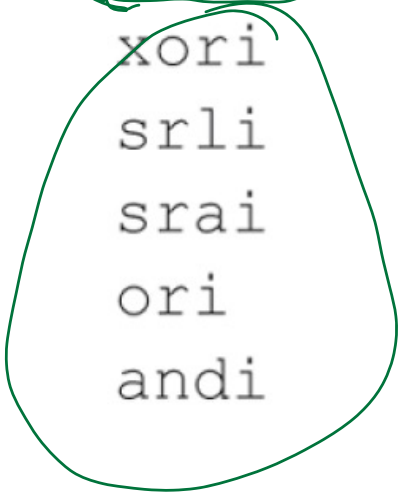
RISC-V I-type Instructions

| imm[11:0] | rs1 | funct3 | rd | opcode |
|-----------|--------|--------|--------|--------|
| 12 bits | 5 bits | 3 bits | 5 bits | 7 bits |

imm[11:0] means bits 11, 10, 9, ..., 0

- Fields in I-type
 - opcode: operation code
 - rd: destination register number
 - rs1: first source register number
 - funct3: additional function code
 - imm: lower 12 bits of the immediate, in the place of funct7 and rs2
- Since only 12 bits are kept in machine code, the immediate must be in $[-2^{11}, +2^{11} - 1]$ or $[-2048, 2047]$

Example of I-type opcode and funct3 code

| | type | opcode | funct3 |
|---|------|---------|--------|
|  | I | 0010011 | 000 |
| slli | I | 0010011 | 001 |
| slti | I | 0010011 | 010 |
| sltiu | I | 0010011 | 011 |
|  | I | 0010011 | 100 |
| xori | I | 0010011 | 101 |
| srli | I | 0010011 | 101 |
| srai | I | 0010011 | 101 |
| ori | I | 0010011 | 110 |
| andi | I | 0010011 | 111 |

In **slli**, **srli**, and **srai**, only lower 5 bits of the immediate are used for shift amount.
5 bits are enough for 32-bit registers!

I-type Example: ADDI



addi x1, x2, 32

rd

addi rd, rs1, imm



I-type Example: ADDI

| imm[11:0] | rs1 | funct3 | rd | opcode |
|-----------|--------|--------|--------|--------|
| 12 bits | 5 bits | 3 bits | 5 bits | 7 bits |

addi x1, x2, 32

| | | | | |
|----------------|-------|-----|-------|---------|
| 0000 0010 0000 | 00010 | 000 | 00001 | 0010011 |
|----------------|-------|-----|-------|---------|

32 bit $\times 2$
12 bit + 32
mismatch

When executing the instruction, processor builds a 32-bit immediate **imm** ?

| | | |
|-----|--------------------------|----------------|
| | imm[31:12] | imm[11:0] |
| imm | 0000 0000 0000 0000 0000 | 0000 0010 0000 |

Higher 20 bits must be the same as the sign

I-Type: shift instructions(SLLI, SRLI, and SRAI)

| imm[11:0] | rs1 | funct3 | rd | opcode |
|-----------|--------|--------|--------|--------|
| 12 bits | 5 bits | 3 bits | 5 bits | 7 bits |

slli rd, rs1, imm

srli rd, rs1, imm

srai rd, rs1, imm

| Instruction | funct3 |
|-------------|--------|
| SLLI | 001 |
| SRLI | 101 |
| SRAI | 101 |

imm is less than 32 in these instructions

We do not need 12 bits. Higher 7 bits are not used. We still call them funct7

The format is still **I-Type**

| funct7 | imm[4:0] | rs1 | funct3 | rd | 0010011 |
|---------|----------|--------|--------|--------|---------|
| 7 bits | 5 bits | 5 bits | 3 bits | 5 bits | 7 bits |
| 0000000 | | | | | |

I-type Example: SRLI vs SRAI

| funct7 | imm[4:0] | rs1 | funct3 | rd | opcode |
|--------|----------|--------|--------|--------|--------|
| 7 bits | 5 bits | 5 bits | 3 bits | 5 bits | 7 bits |

srli x1, x2, 16

| | | | | | |
|---------|-------|-------|-----|-------|---------|
| 0000000 | 10000 | 00010 | 101 | 00001 | 0010011 |
|---------|-------|-------|-----|-------|---------|



Same opcode and funct3
Bit 30 is different!

srai x1, x2, 16

| | | | | | |
|---------|-------|-------|-----|-------|---------|
| 0100000 | 10000 | 00010 | 101 | 00001 | 0010011 |
|---------|-------|-------|-----|-------|---------|



| | | | | | |
|------|---|---------|-----|---------|----------|
| srli | I | 0010011 | 101 | 0000000 | ← funct7 |
| srai | I | 0010011 | 101 | 0100000 | |

What format would you use for load instructions?

- Load instructions

`lw rd, offset(rs1)`

A. R-format

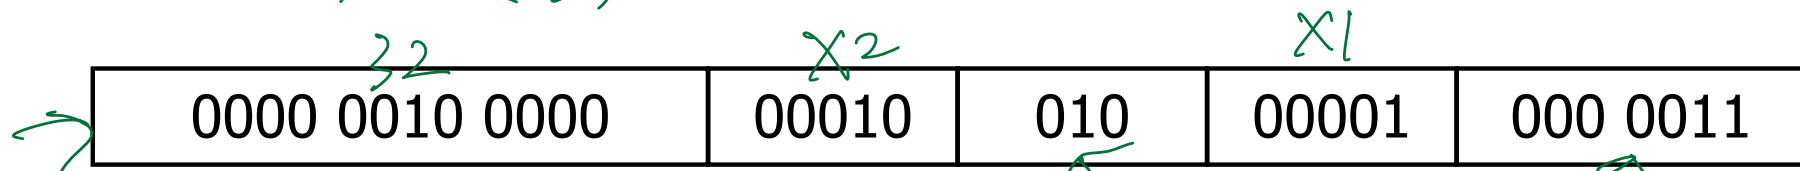
☒ B. I-format

| | | | | | | |
|------|--------------------------|-----|-----|--------|----|--------|
| → R: | funct7 | rs2 | rs1 | funct3 | rd | opcode |
| → I: | imm [11:0] ✓ offset ✓ | | rs1 | funct3 | rd | opcode |

Loads are I-type



lw x1, 32(x2)
lw Rd, imm(RS1)



| | type | opcode | funct3 |
|------|------|---------|--------|
| lb | I | 0000011 | 000 |
| lh | I | 0000011 | 001 |
| → lw | I | 0000011 | 010 |
| ld | I | 0000011 | 011 |
| lbu | I | 0000011 | 100 |
| lhu | I | 0000011 | 101 |
| lwu | I | 0000011 | 110 |

Do you see the pattern in funct3?

How about store instructions?

- Store instructions have rs2, but not rd

~~sw~~ ~~rs2,offset(rs1)~~

| | | | | | | |
|----|------------|-----|-----|--------|---------------|--------|
| R: | funct7 | rs2 | rs1 | funct3 | rd | opcode |
| I: | imm [11:0] | | rs1 | funct3 | rd | opcode |

need a new format!

How to design?

RISC-V S-type Instructions



sw rs2,offset(rs1) # imm is offset

- Fields in S-type

- opcode: operation code
- rs1: first source register number
- rs2: second source register number
- imm[11:5] and imm[4:0]:

The lower 12 bits of the immediate are stored into two fields

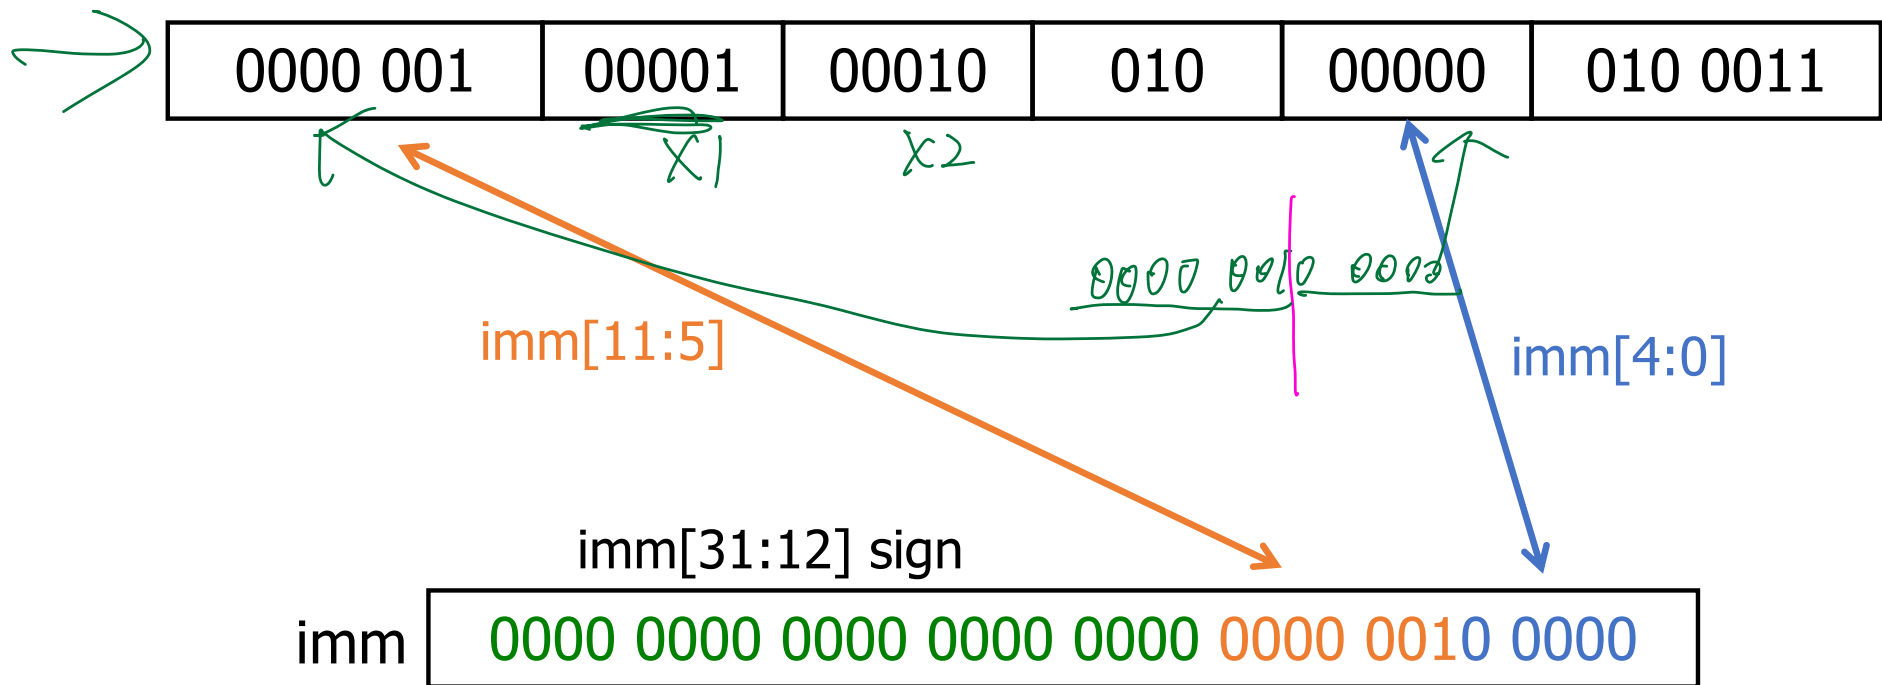
Bits 11 to 5, are in funct7 and bits 4 to 0, are in rd

Stores are S-type



sw x1, 32(x2)

sw rs2, imm(rs1)
(X1) (X2)



imm[11:0] are saved in two fields

Summary of R-, I-, and S-type instructions

| Instruction | Format | funct7 | rs2 | rs1 | funct3 | rd | opcode |
|-------------|--------|---------|-----|-----|--------|-----|---------|
| add (add) | R | 0000000 | reg | reg | 000 | reg | 0110011 |
| sub (sub) | R | 0100000 | reg | reg | 000 | reg | 0110011 |

| Instruction | Format | immediate | rs1 | funct3 | rd | opcode |
|----------------------|--------|-----------|-----|--------|-----|---------|
| addi (add immediate) | I | constant | reg | 000 | reg | 0010011 |
| lw (load word) | I | address | reg | 010 | reg | 0000011 |

| Instruction | Format | immed-iate | rs2 | rs1 | funct3 | immed-iate | opcode |
|-----------------|--------|------------|-----|-----|--------|------------|---------|
| sw (store word) | S | address | reg | reg | 010 | address | 0100011 |

Fields, other than immediate fields, are located at the same location, for all types
 Placement of bits in the immediate is more complicated

RISC-V Core Instruction Format

- Now we know three types: R, I, and S
 - We will discuss other formats later

| | 31 | 25 | 24 | 20 | 19 | 15 | 14 | 12 | 11 | 7 | 6 | 0 | | |
|-----------|-----------------------|----|----|----|-----|-----|-----|--------|--------|----|-------------|--------|--------|--|
| R | funct7 | | | | rs2 | | rs1 | | funct3 | | rd | | opcode | |
| I | imm[11:0] | | | | | rs1 | | funct3 | | rd | | opcode | | |
| S | imm[11:5] | | | | rs2 | | rs1 | | funct3 | | imm[4:0] | | opcode | |
| SB | imm[12 10:5] | | | | rs2 | | rs1 | | funct3 | | imm[4:1 11] | | opcode | |
| U | imm[31:12] | | | | | | | | | rd | | opcode | | |
| UJ | imm[20 10:1 11 19:12] | | | | | | | | | rd | | opcode | | |

We may use funct7, rs2, and rd to refer to a group of bits, although some formats do not have these fields. For example rs2 always bits 20 to 24, even if I-type instructions do not have rs2 field.

Pseudoinstructions are converted to real instructions first, and then to machine code.

Study the remaining slides yourself

Examples of R-, I-, and S-type instructions

| R-type Instructions | funct7 | rs2 | rs1 | funct3 | rd | opcode | Example |
|----------------------|--------------|-------|-------|--------|------------|---------|-------------------|
| add (add) | 0000000 | 00011 | 00010 | 000 | 00001 | 0110011 | add x1, x2, x3 |
| sub (sub) | 0100000 | 00011 | 00010 | 000 | 00001 | 0110011 | sub x1, x2, x3 |
| I-type Instructions | immediate | | rs1 | funct3 | rd | opcode | Example |
| addi (add immediate) | 001111101000 | | 00010 | 000 | 00001 | 0010011 | addi x1, x2, 1000 |
| lw (load word) | 001111101000 | | 00010 | 010 | 00001 | 0000011 | lw x1, 1000 (x2) |
| S-type Instructions | immed-iate | rs2 | rs1 | funct3 | immed-iate | opcode | Example |
| sw (store word) | 0011111 | 00001 | 00010 | 010 | 01000 | 0100011 | sw x1, 1000(x2) |

Exercise

- Pick an instruction and encode it
- Study the machine code generated by RARS

Example

Hex: 00030503

Bin: 0000 0000 0000 0011 0000 0101 0000 0011

Fields: 0000 0000 0000 0011 0000 0101 0000 0011

| | | | |
|-----|---|---------|-----|
| lb | I | 0000011 | 000 |
| lh | I | 0000011 | 001 |
| lw | I | 0000011 | 010 |
| ld | I | 0000011 | 011 |
| lbu | I | 0000011 | 100 |
| lhu | I | 0000011 | 101 |
| lwu | I | 0000011 | 110 |

Questions

- RISC-V has three fields for specifying operations: opcode, funct3, and funct7. Why don't they combine them and have a single opcode field of 17 bits?
- If someone decides to increase the number of registers to 64, how does it affect the encoding of instructions?