

Functions



Z. Jerry Shi

Department of Computer Science and Engineering
University of Connecticut

CSE3666: Introduction to Computer Architecture

Outline

- Function basics
 - Write a function
 - Call a function
 - Calling convention
- Using stack to save/restore data

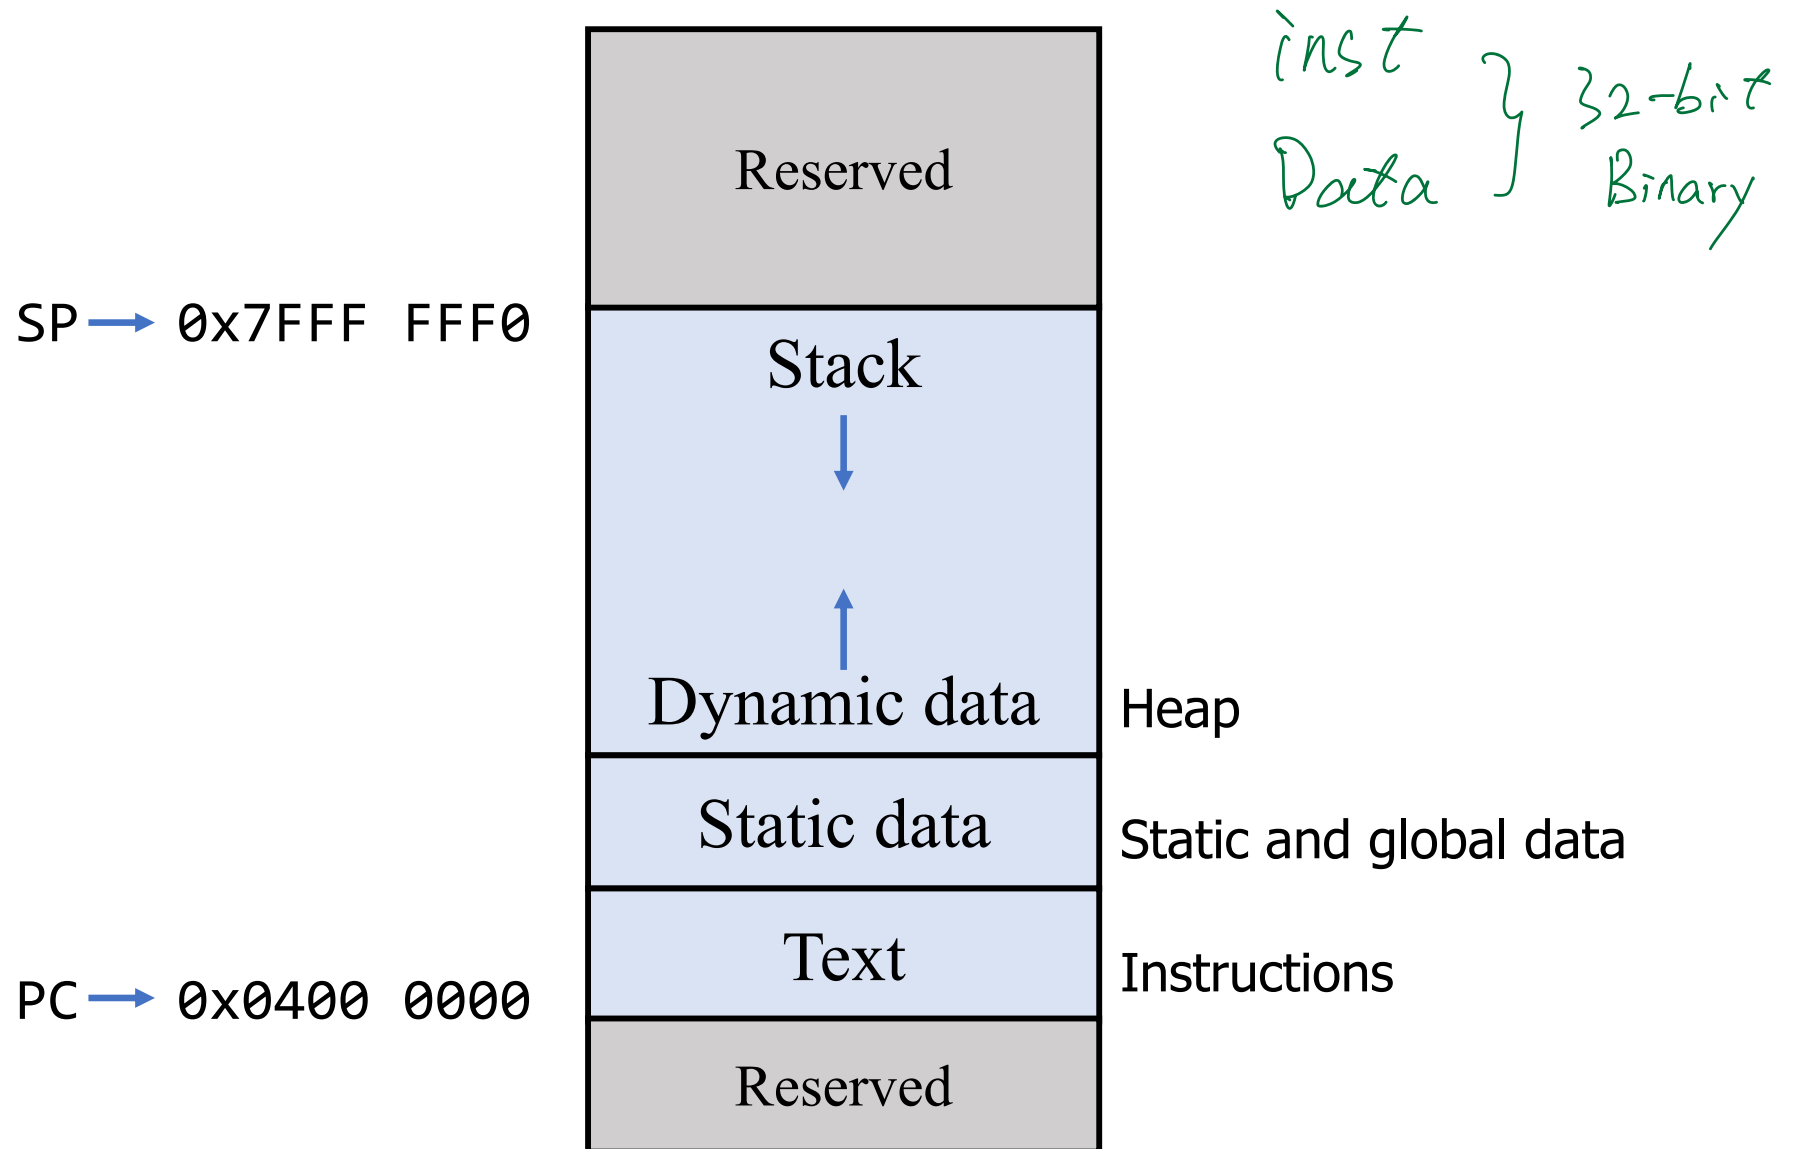
Reading: Sections 2.8

References: Reference card in the book

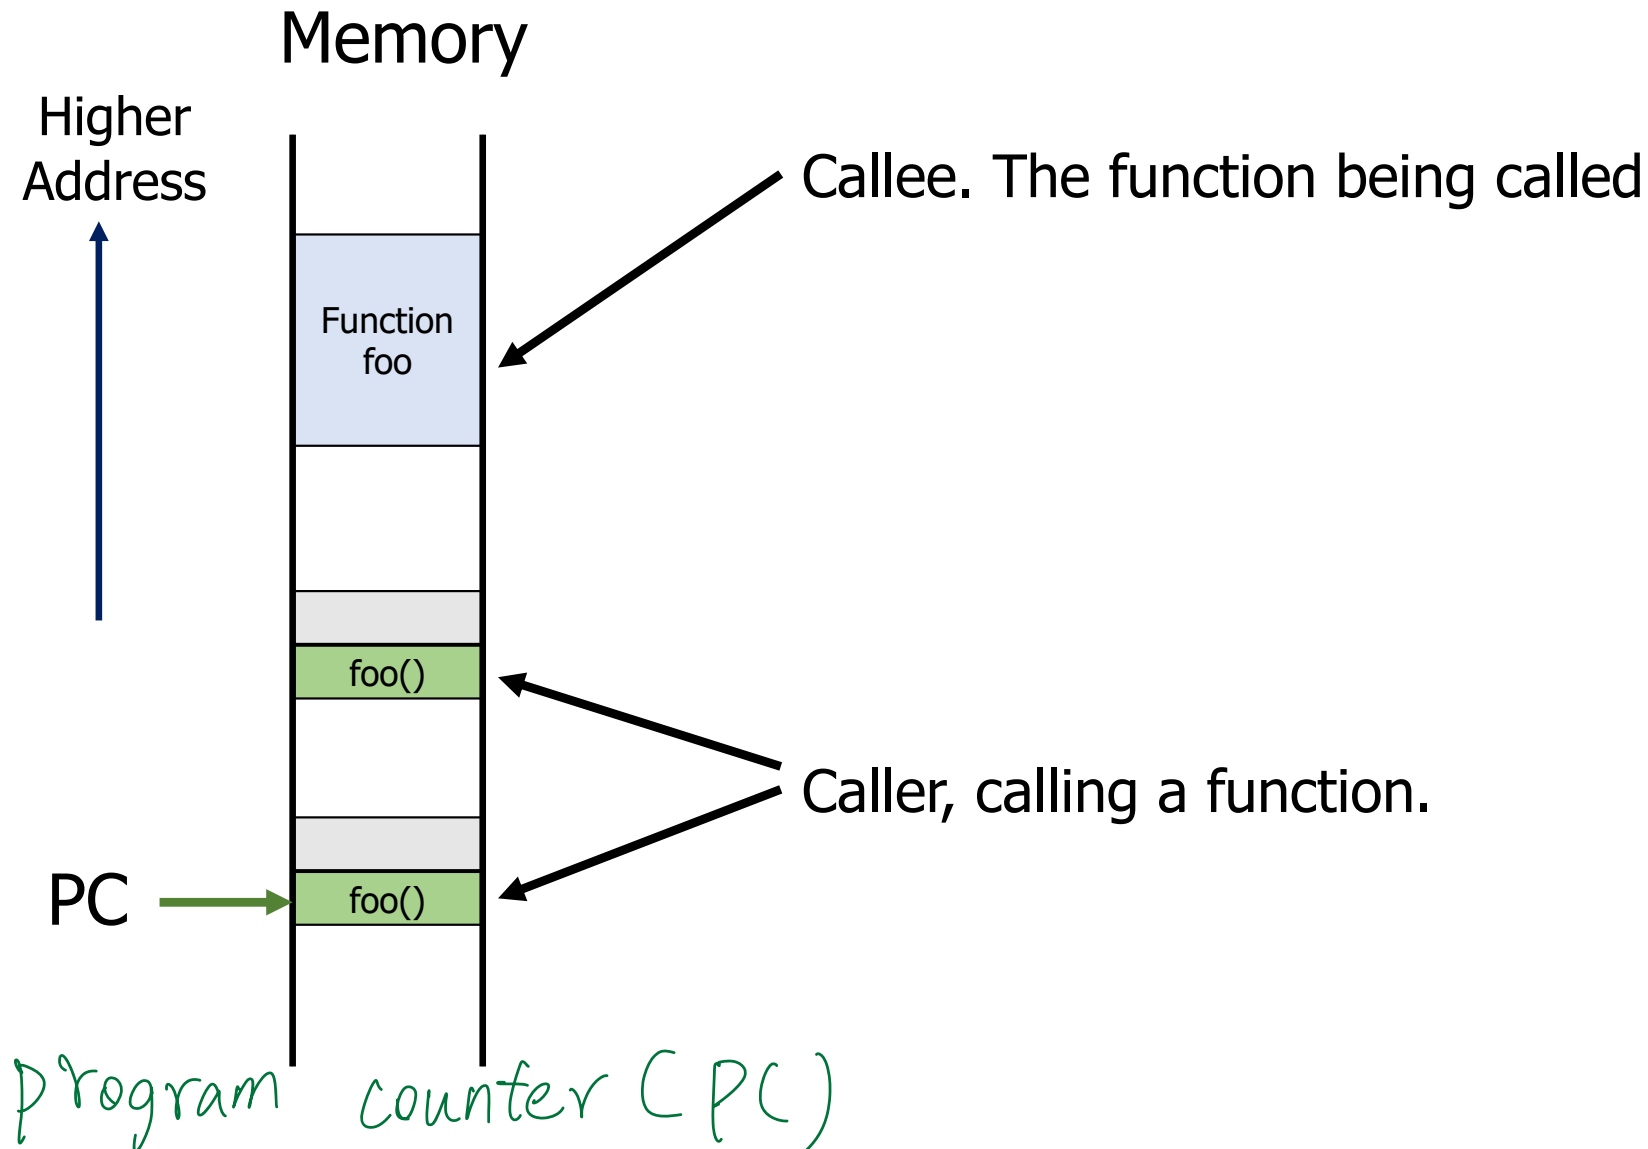
Functions

- Make code manageable
 - Hierarchical design
- Reduce code size
 - You do not have to write the same code again and again
- Easy to maintain
 - No need to find all the copies of the code to fix a bug
 - New versions of functions are called if the library is dynamically linked
- ...

Memory Layout

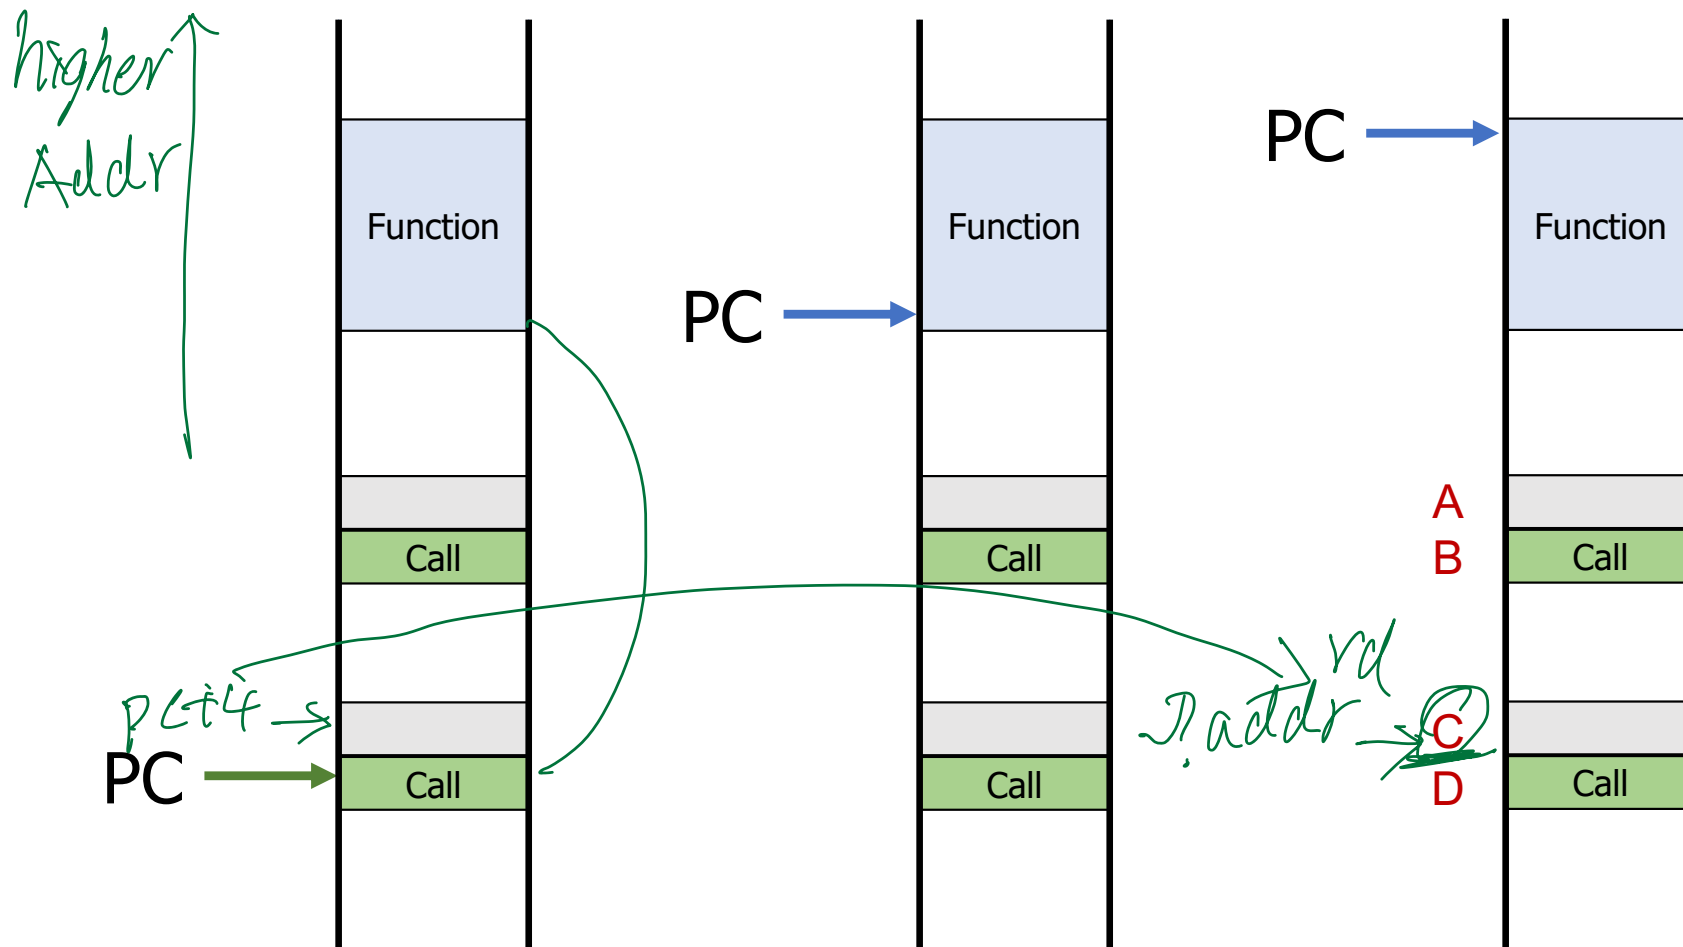


Flow with functions




Flow with functions

- Upon function call, processor starts to execute instructions in functions
- When function returns, **where should the processor get the next instruction?**
 - Select one from four memory locations: A, B, C and D.



Flow with functions

- Processor **needs to know the return address** to resume flow
 - Branches won't work!
- Processor should resume from **the instruction that follows the function call**
 - We need to keep the information at function call



RISC-V Procedure Call Instructions - JAL

`jal rd, immdfunction_name` ← Function name is just a label, indicating a memory address

- jump and link (JAL) for procedure call

① – Go to the label (where the procedure/function starts)

② ~~XX~~ And save **the address of the following instruction (PC+4)** in **rd**

↓
next inst

- rd can be any register ^{↗ return addr}
 - **Save return address in ra (or x1) for function calls**
 - If the return address is not needed, use x0 as rd

- J is a pseudoinstruction for jump

- It is a JAL instruction that writes to x0
- There is no separate jump instruction in RISC-V

How do we return?

- Now the return address is in a register, e.g., ra. How can we actually go there?
 - JAL does not take a source register
- We need another instruction that can jump to the address in a register

JALR

JALR

`jalr rd, rs1, offset`

- Procedure call or return: **jump and link to address in a register**
 - Go to the address computed as $\text{Reg}[\text{rs1}] + \text{offset}$
 - Saves **the address of the following instruction (PC+4)** in **rd**
- The return address is saved in **rd**
 - If the return address is not needed, use **x0** as **rd**
 - **When used for function calls, save the return address in **ra** (or **x1**)**
 - Yes, JALR can be used to call functions

In this course, it is most commonly used for return

`jalr x0, ra, 0`

Function Call Examples

main:

jal ra,foo

add s0,s0,a0

1. JAL sets PC = foo, ra = main + 4

foo:

add t1,a0,x0

.....

xor t0,t2,t3

.....

jalr x0,ra,0

2. Continue from foo

3. JALR sets PC = ra

main() is the caller. foo() is the callee().

foo() is a **leaf function/procedure**, which does not call other functions

Example

- Write a function to calculate the absolute value of a word

```
int    abs (int    n)
{
    int    rv = n;
    if (n < 0)
        rv = -n;
    return rv;
}
```

How does the function get the argument?

How does the function pass the return value to the caller?

Passing parameters and returning values

- First eight parameters/arguments are passed to the function in

a0, a1, a2, a3, a4, a5, a6, a7



- Two values can be returned in
 - If only one value is returned, it is in a0

a0, a1

Caller and callee communicate through these registers only
Not through registers like s1, s2

For example, a function cannot assume s1 has an address it needs

Now, we can continue

Label the entry point



```
# RISC-V
abs:
```

```
int abs (int n)
{
    int rv = n;
    if (n < 0)
        rv = -n;
    return rv;
}
```

Translate the first statement:

n is in a0

Which register do we assign to rv?

abs function

abs: # when the function is called, a0 is n

bge a0, x0, exit # n < 0?

sub a0, x0, a0 # rv = -n

exit:

jalr x0, ra, 0 # return

The function has 3 instructions

example of calling abs(-2)


this is NOT part of function implementation

addi a0, x0, -2 # set a0

jal ra, abs # call abs

a0 should be 2 here

Nested Function Call Examples



Do you see any problems?

The problems

- Problem 1: In foo, register ra is changed when calling bar
- Problem 2: foo does not know bar changed register t1


How do we deal with the problems?

Stack

- Stack grows **from higher address toward lower address**
 - **sp** is the address of the word at the top of the stack
- Two operations:
 - **push** adds data to the top of the stack
 - **pop** removes data from the top of the stack
- What can stack be used for?
 - Saving registers
 - Keeping local variables (used by a function)
 - Passing Arguments
 - Returning values

*First-in
last-out*

Stack operations



#push x1

① `addi sp, sp, -4`


② `sw x1, 0(sp)`

pop x1

① `lw x1, 0(sp)`


② `addi sp, sp, 4`

Memory




Push

Memory



Memory



Exercises

- Write RISC-V instructions for the following operations

Push ra onto the stack

*addi sp, sp, -4
sw ra, 0(sp)*

Pop ra from the stack

Push s1, s2, and s3 onto the stack

How many instructions do you need for each of them?

Example: push multiple words

- Write RISC-V instructions to push s1, s2, and s3 onto the stack

Instead of 3 ADDI's, we can adjust sp once for three words

reserve space for 3 words

```
addi sp, sp, -12  
sw    s1, 8(sp)  
sw    s2, 4(sp)  
sw    s3, 0(sp)
```

Before push

sp →

After push

sp →

Address	Value
0x7FFF 901C	X
0x7FFF 9018	X
0x7FFF 9014	X
0x7FFF 9010	s1
0x7FFF 900C	s2
0x7FFF 9008	s3
0x7FFF 9004	
0x7FFF 9000	

Push and pop pairs

push s1, s2, and s3

addi sp, sp, -12

sw s1, 8(sp)

sw s2, 4(sp)

sw s3, 0(sp)

pop s1, s2, and s3


lw s1, 8(sp)

lw s2, 4(sp)

lw s3, 0(sp)

addi sp, sp, 12

restore sp after loads



Address	Value
0x7FFF 901C	
0x7FFF 9018	
0x7FFF 9014	
0x7FFF 9010	s1
0x7FFF 900C	s2
0x7FFF 9008	s3
0x7FFF 9004	
0x7FFF 9000	

Nested function example: fact

- C code:

```
int fact(int n)
{
    if (n < 1)
        return 1;
    else
        return n * fact(n - 1);
}
```

Argument n in a0

Result in a0

Fact – Will this work?

fact:

```
# if (n < 1) return 1
addi t0, zero, 1
bge a0, t0, ELSE    # go to else branch if n >= 1
```

```
# if branch
addi a0, zero, 1    # set return value 1
beq  x0, x0, EXIT    # go to exit
```

```
ELSE:    # else n * fact(n - 1)
addi a0, a0, -1    # compute fact(n-1)
jal  ra, fact
mul  a0, a0, a0    # n * (n - 1)!
```

```
EXIT:
jalr  x0, ra, 0    # Use only one exit !
```


RISC-V code for fact

```
fact:  # One entrance and one exit
      addi sp, sp, -8      # Save return address and n
      sw    ra, 4(sp)
      sw    a0, 0(sp)
      addi t0, zero, 1
      bge a0, t0, ELSE    # if n >= 1, go to else branch
      addi a0, zero, 1    # return 1
      beq x0, x0, EXIT    # goto exit

ELSE:
      addi a0, a0, -1      # call with n - 1
      jal  ra, fact
      lw    t0, 0(sp)      # restore original n
      mul a0, a0, t0        # n * (n - 1)!
EXIT:  lw    ra, 4(sp)      # restore return address
      addi sp, sp, 8        # restore sp
      jalr x0, ra, 0        # and return
```

Question

- We double check and find regs that are not preserved in the function
- We changed a0 and t0. Do we need to save them?

What registers are preserved through function call?

Does the callee need to save/restore every register?

- It needs registers to do its job

If F calls G, F is the caller and G is the callee

What registers can be used by callee?

Callee does not need to preserve t0-t6 and a0-a7

Callee can use them without saving/restoring

Caller cannot assume these registers are preserved

Callee needs to preserve other registers

Caller assumes these registers are preserved

Preserved	Not preserved
Saved registers: x8-x9, x18-x27	Temporary registers: x5-x7, x28-x31
Stack pointer register: x2(sp)	Argument/result registers: x10-x17
Frame pointer: x8(fp)	
Return address: x1(ra)	
Stack above the stack pointer	Stack below the stack pointer

How does callee preserve a register?

Question

- What kind of registers is better to keep *i* in function *foo*?

A. temporary

B. argument

C. saved

```
int    foo()  
{
```

```
    int i;
```

```
    // some instructions
```

```
    for (i = 0; i < 10; i ++)  
        bar(i);
```

```
    // more instructions
```

```
}
```

to

loop

push i

call bar

pop i

so

push *SO*

loop

call bar

Local storage

- Storage for local variables is allocated on the stack

```
int foo()  
{  
    // keep both a and b on stack  
    int a[15], b;  
    ...
```

```
foo:  
# reserve space for 16 words  
addi    sp, sp, -64
```

If needed, allocate more space for other local variables and saved registers

*do ~ a₁₄? ✓
do ~ a₁₅?*

+4

	Address	Value
	0x7FFF 9040	
b	0x7FFF 903c	<i>b</i>
	0x7FFF 9038	<i>a[14]</i>
	...	
a	0x7FFF 900c	
	0x7FFF 9008	
	0x7FFF 9004	
sp	0x7FFF 9000	<i>a[0]</i>

Passing arrays or strings to a function

- The starting address of the array is passed to functions
 - Changes to the array are preserved after the function returns
 - An ASCII string is an array of bytes, ending with a null character

```
int  array_max (int arr[], int length);  
unsigned int strlen (char s[]);
```

The caller puts the starting address of arr or s in a0.

Example of passing a string to a function

```
unsigned int strlen (char s[]);
```

The caller puts the correct address in a0.

For example,

strlen(s)

```
la    a0, s
# a0 = 0x00FE9000
jal   ra, strlen
# a0 = _____
```

Address	Value
0x00FE 9007	0
0x00FE 9006	54 G
0x00FE 9005	54 6
0x00FE 9004	54 6
0x00FE 9003	51 3
0x00FE 9002	69 E
0x00FE 9001	83 S
0x00FE 9000	67 C

Example: strlen

- C function strlen() returns length of a (null-terminated) string

```
unsigned int strlen (char s[])
```

```
{  
    unsigned int i;  
    i = 0;  
    while (s[i] != 0)  
        i += 1;  
    return i;  
}
```

strlen() does not call other functions.
It is a **leaf procedure**.

- When implementing the function with RISC-V,
 - Where do we find the address of the string?
 - Where do we keep i? a, t, or s registers?

RISC-V strlen

```
strlen:      # a0 is the starting address
             add    t1, x0, x0          # i = 0
loop:
             add    t0, t1, a0          # addr of s[i]
             lb     t0, 0(t0)           # load a byte
             beq    t0, x0, exit
             addi   t1, t1, 1           # i += 1
             beq    x0, x0, loop        # goto loop
exit:
             addi   a0, t1, 0           # set return value
             jalr   x0, ra, 0           # return
```

Question

- Which register is used to pass parameters to a function?
- A. a0
- B. t0
- C. s0
- D. sp

Question

- If a function returns a value, where can the caller find it?
- A. `a0`
- B. `t0`
- C. `s0`
- D. `sp`

Question

- Which register should a function save before changing it?
- A. a0
- B. t0
- C. s0
- D. sp

Question

- When implementing the following function with RISC-V assembly language, in which register is argument c stored at the entry of the function?

```
int  foo(int  a[], int b, char c[], int d);
```

- A. a0
- B. a1
- C. a2
- D. a3

RISC-V Calling Convention

- Parameter passing
 - The first eight arguments are in a0, a1, a2, a3, a4, a5, a6, and a7
 - More arguments are placed on stack
- Register preservation and restore
- Return values
 - Return values are placed in a0 and a1

Assume argument registers (like a0 and a1) and temporary registers (like t0 and t1) are changed during any function call although they are not in some cases

RISC-V Calling Convention

[riscv-elf-psabi-doc/riscv-cc.adoc at master · riscv/riscv-elf-psabi-doc \(github.com\)](https://github.com/riscv/riscv-elf-psabi-doc)

Study the remaining slides yourself

Alignment on stack

- RISC-V calling convention requires the stack pointer is always 16-byte aligned.
 - Even if you need space for one word, compiler reserves 16 bytes
- In this course, we keep the stack pointer aligned at a word address
 - 4-byte aligned

Study the sort example in textbook

- Study the code and answer the following questions
 - Which registers are used to store loop control variables i and j? Is it better to assign temporary registers (like t0 and t1) to i and j?
 - What registers are saved and restored in sort function?
- Write the sort function without looking at the assembly code provided in the textbook

RISC-V register conventions

Name	Register number	Usage	Preserved on call?
x0	0	The constant value 0	n.a.
x1 (ra)	1	Return address (link register)	yes
x2 (sp)	2	Stack pointer	yes
x3 (gp)	3	Global pointer	yes
x4 (tp)	4	Thread pointer	yes
x5 - x7	5–7	Temporaries	no
x8 - x9	8–9	Saved	yes
x10 - x17	10–17	Arguments/results	no
x18 - x27	18–27	Saved	yes
x28 - x31	28–31	Temporaries	no

General Procedure Calling Steps

Caller

1. Save (t, a, and ra) registers, if needed
2. Set parameters
3. Transfer control to function (JAL, JALR)

Callee

1. Acquire storage for local variables/register saving
2. Save registers, if necessary
3. Do the work
4. Set return value
5. Restore registers and stack
6. Return (JALR)

4. Restore registers and stack

Tips

- Write function body first, and then check what registers need to be saved
- Save registers at the beginning of a function
- Restore registers and stack just before returning
 - Use one exit from the function. **Only one return!**

```
add    a0, s1, s2    # return s1 + s2
beq    x0, x0, exit  # goto exit
...
li     a0, -1        # return -1
```

exit:

```
# restore registers and then return
```

Pitfalls

- Forget calling conventions
 - Some values are not preserved in function calls
 - Assume certain values are in temporary or saved registers (e.g., s1)
 - Place return value in wrong registers
 - Forget to save ra in non-leaf functions
 - Forget to save temporary registers in non-leaf functions
 - Forget to restore sp
- Too many exits from a function
 - Easy to forget to fix bugs at all exits
- Mix up code from different functions

Pitfalls

- Access the space that is already freed

For example, adjust sp first before loading saved values

```
# Wrong order ! Load first  
addi    sp, sp, 4           # space already released!  
lw      ra, -4(sp)
```

A common mistake In C: returning a pointer to local variables

Source code

Function names are just labels in the source code!

Source code

```
main:  
    jal  bar  
    nop  
    .....
```

```
bar:  
    addi t1,...  
    .....
```

```
foo:  
    xor t0,...  
    .....
```

Instructions in functions do
not overlap

Exercise: max3

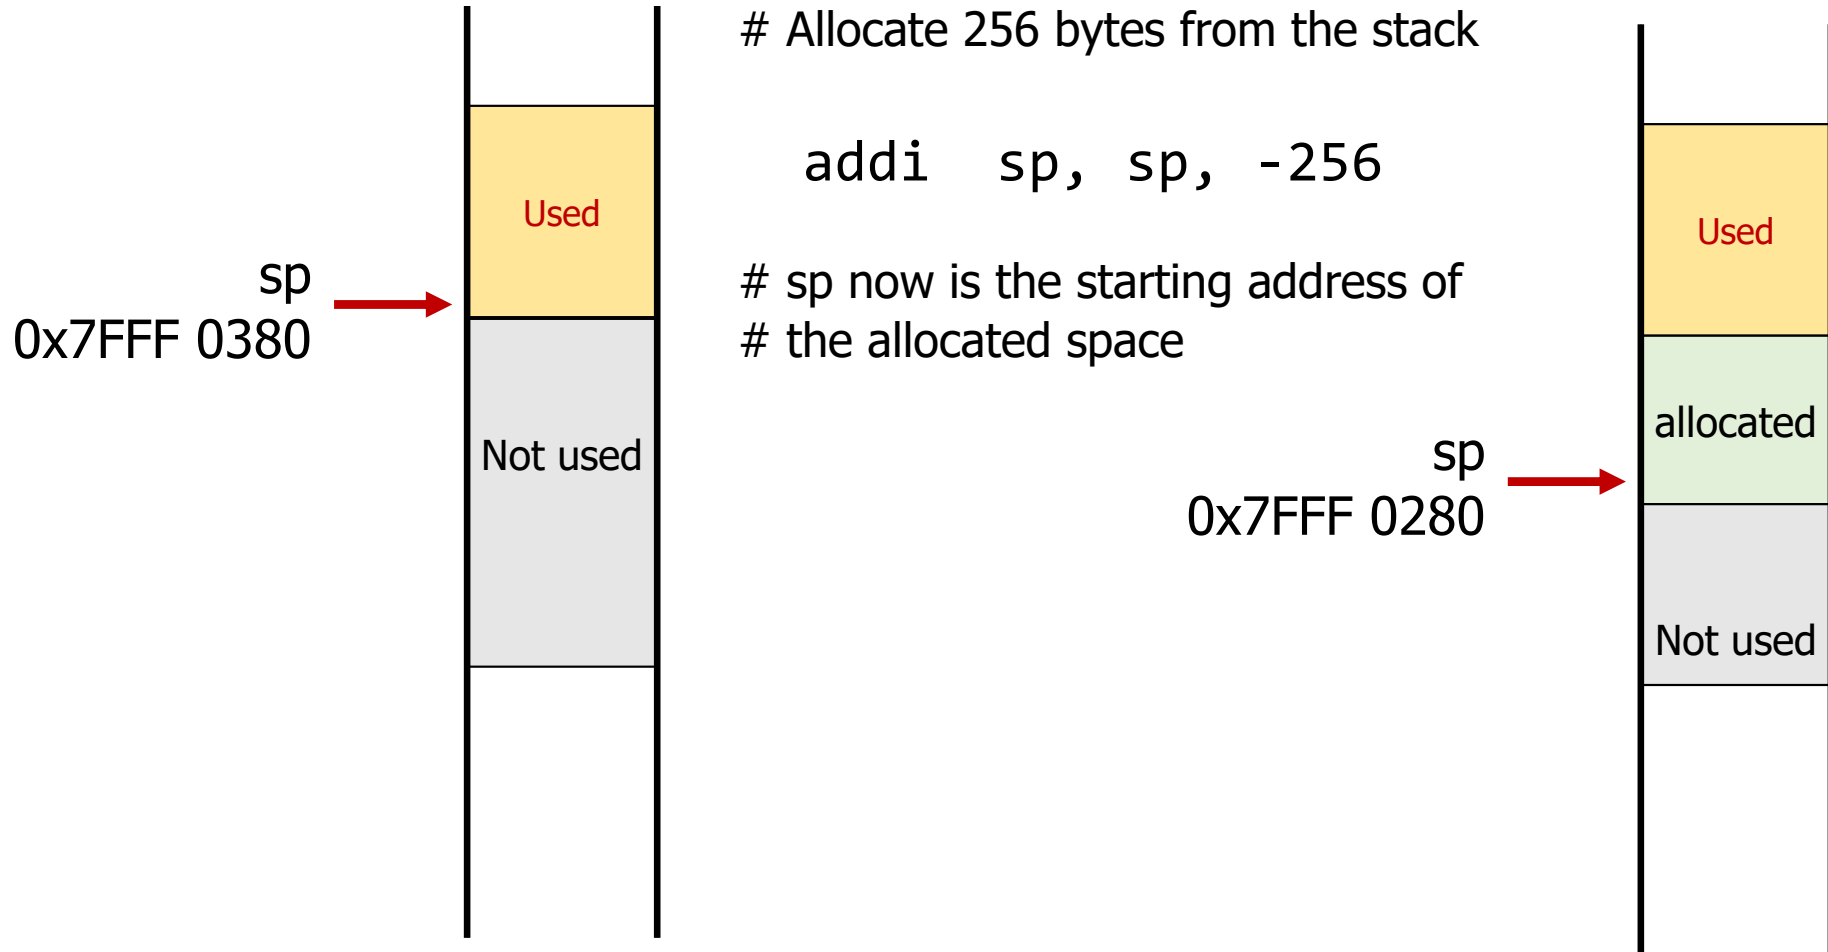
```
int max3 (int a, int b, int c)
{
    int i = a;

    if (i < b)
        i = b;
    if (i < c)
        i = c;
    return i;
}
```


Exercise

- Write code to call `max3()`
`m = max3(10, 20, -20)`
- Write examples in the textbook
 - `sort`
 - `strcpy`
- Write more code
 - For example, a function that finds the median

Allocating space on the stack



Use JALR

- We use JALR to return from a function

```
jalr x0, ra, 0
```

```
# the following two are pseudoinstructions
```

```
jr    ra
```

```
ret
```

- We can call a function that is far away
 - Place the function's address in a register
 - JAL cannot jump too far
- We can also change the function we calls
 - Function pointers, and virtual functions