

Floating-Point Numbers



Z. Jerry Shi

Department of Computer Science and Engineering
University of Connecticut


CSE3666: Introduction to Computer Architecture

Outline

- Real numbers in binary
 - Decimal to binary
 - Binary to decimal
- IEEE 754 floating-point number standards
 - Single precision and double precision
- RISC-V support for floating-point numbers

Reading: **Section 3.5**, excluding hardware support for floating-point numbers.

$0.00\dots001$ $0.999\dots9$



Real numbers

- Computers need to deal with
 - Numbers with fractions (not just whole numbers)
 - Very big numbers
 - Very small numbers

Example of real numbers in decimal:

- ? 3.14159... not normalized
- ? -0.002×10^{-20}
- ? 9.4607×10^{15} (meters in a light year)

Normalized scientific notation:


Only one non-zero digit to the left of the decimal point.

Binary number with fraction

exact match

- To represent fractions in binary, we use bits after the binary point

What is the value of the following binary number?



*radix ?
(base) :*

$$2^2 \ 2^1 \ 2^0 \ . \ 2^{-1} \ 2^{-2} \ 2^{-3}$$

Binary to decimal

Example: $0b101.1101 = 5.8125_{(10)}$

bits	1	0	1	1	1	0	1
weights	2^2	2^1	2^0	2^{-1}	2^{-2}	2^{-3}	2^{-4}

(base radix)

Multiply each bit with weight:

$$\begin{aligned} &0b101.1101 \\ &= 1 \times 2^2 + 0 \times 2^1 + 1 \times 2^0 \\ &\quad + 1 \times 2^{-1} + 1 \times 2^{-2} + 0 \times 2^{-3} + 1 \times 2^{-4} \\ &= 4 + 0 + 1 + 0.5 + 0.25 + 0 + 0.0625 \\ &= 5.8125 \\ &\quad [Dec] \end{aligned}$$

Integer part

Fractional part

Decimal to binary

Example:

Convert the decimal number 0.8 to a binary number

0						
2^0	2^{-1}	2^{-2}	2^{-3}	2^{-4}	2^{-5}	2^{-6}

Converting decimal to binary

$$0.8 = 0.11001100\dots$$

Decimal	Binary
0.8	0.
0.8 * 2 = 1.6	0.1
0.6 * 2 = 1.2	0.11
0.2 * 2 = 0.4	0.110
0.4 * 2 = 0.8	0.1100
0.8 * 2 = 1.6	0.11001...
Continue....	0.11001100110011001100 ...

Fraction .8 appears again. The pattern 1100 will repeat forever.

python

```
>>> float.hex(0.8)
'0x1.999999999999ap-1'
```

$$1.100110011001\dots \times 2^0 \quad \text{bin}$$

$$1.9999\dots \times 2^0 \quad \text{hex}$$

Normalized notation of binary numbers

- There are many representations as we move the binary point

$$101.1101 = 10.11101 \times 2^1 = \boxed{1.011101 \times 2^2} = 0.1011101 \times 2^3$$

Normalized binary representation

The **normalized binary representation** has a single 1 before the point

$$\pm 1.x \times 2^E$$

Significand

Exponent
is written in decimal for
convenience

python

```
>>> float.hex(float.fromhex('5.d'))  
'0x1.74000000000000p+2'
```


Encode floating-point numbers

- Given a number of bits, how do we represent

Design Problem ? $\pm 1.x \times 2^E$

- What need to be encoded? \pm, x, E
- How many bits for each?

$$1 \text{ bit} + 29 + 2 = 32 \text{ bits} \quad \text{FP}$$

$$\begin{array}{r} 26 \quad 5 \\ \hline 23 \quad 8 \end{array}$$

IEEE FP

FP8 (meta) ML 9

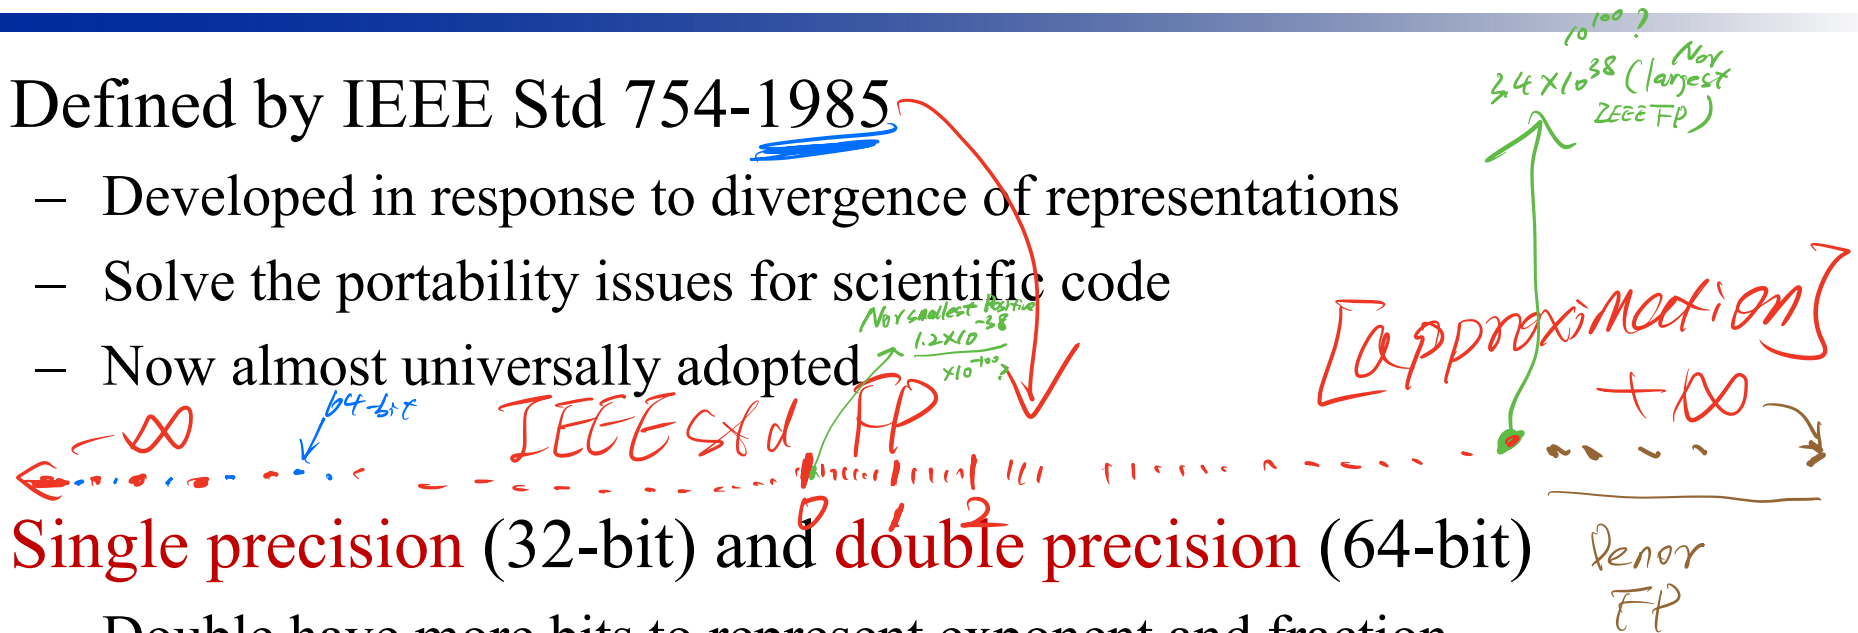
$-\infty$

Real numbers


$+\infty$

Floating Point Standard (single and double precisions)

- Defined by IEEE Std 754-1985
 - Developed in response to divergence of representations
 - Solve the portability issues for scientific code
 - Now almost universally adopted
- Single precision (32-bit) and double precision (64-bit)
 - Double have more bits to represent exponent and fraction
 - They are types float and double in C
- Later versions of the standard include more types
 - E.g., 128-bit quad-precision



IEEE Floating-Point Format: single-precision



$$\text{value} = (-1)^S \times (1.\text{Fraction}) \times 2^E$$

Exponent is in **excess-127 representation**. The Bias = 127.

? ① $\text{ActEx} = \text{EncEx} - 127$

? ② $\text{EncodedExponent} = \text{ActualExponent} + 127$

Exponent field in single-precision

- The exponent field has 8-bit, keeping a value in [0, 255]

- [1, 254]: A **normal** SP number

- We will discuss 0 and 255 soon = *denormal*


- The range of actual exponent: [-126, 127] : *Act Exp*

- Excess-127 representation!

$$\pm 1.x \times 2^E \quad \text{and} \quad E \in [-126, 127]$$

$$\text{Encoded} = E + 127$$

Bits in the exponent field
1 .. 254



Questions: Excess-127

- Given the eight bits in the ^{Exp} exponent field of single-precision FP numbers, find the actual exponents in decimal.

0111 1111 -127

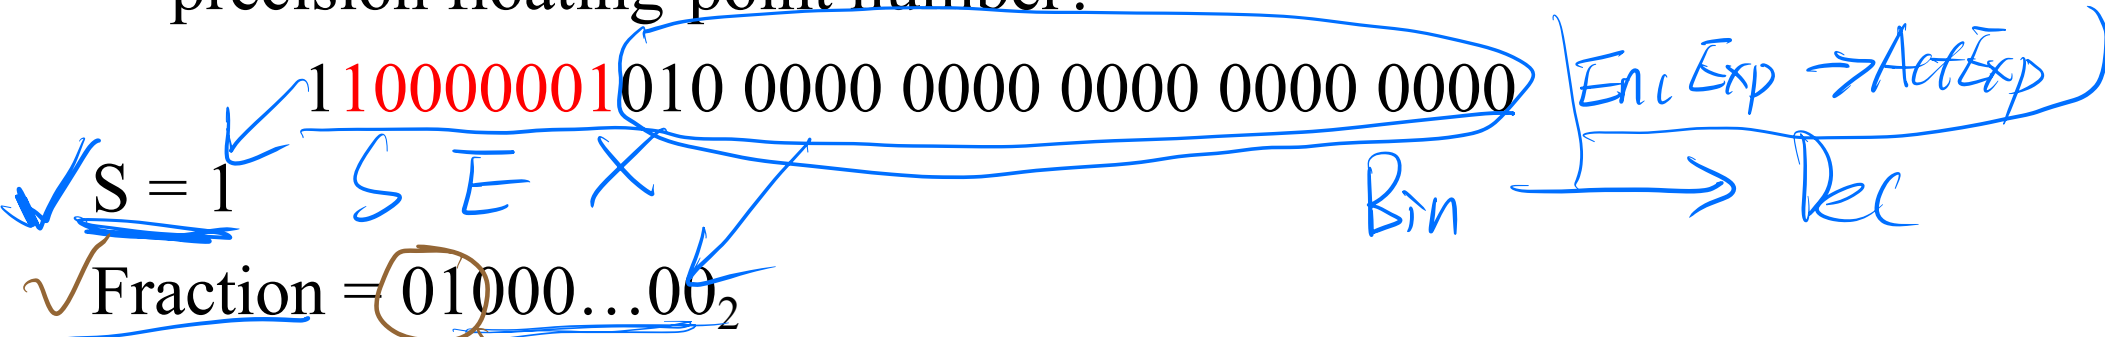
0000 0100 -127

1000 0001 -127

1001 0000 -127

Example: Read Single-Precision FP numbers

- What number (in decimal) is represented by the following single-precision floating-point number?



Encoded exponent = $10000001_2 = 129$ (as 8-bit unsigned number)

Actual exponent = $129 - 127 = 2$

The value is

$$\begin{aligned} & (-1)^1 \times (1 + \underline{0.01}_2) \times 2^{(129 - 127)} \\ &= (-1) \times 1.25 \times 2^2 \\ &= -5 \end{aligned}$$


$Dec \rightarrow Bin$
 $Act\ Exp \rightarrow Enc\ Exp$

Question

What is the actual exponent of the following single-precision floating-point number?

What is its value in decimal?

0x C1C0 0000



Example: Convert to Single-Precision FP numbers

Represent 4.75 with a single precision floating-point number

↓

$$\pm 1.x \times 2^{\text{ActExp}}$$

↓

$$\pm 1.x \times 2^{\text{EncExp} - 127}$$

↓

$$\begin{array}{|c|c|c|} \hline S & \text{EncExp} & \text{Frac} \\ \hline 1 & 8\text{ bit} & X \\ & & (23\text{-bit}) \\ \hline \end{array}$$

Solutions

Represent 4.75 with a single precision floating-point number

$$4.75 = 100.11_2 = (-1)^0 \times 1.0011_2 \times 2^2$$

$$S = 0$$

$$\text{Fraction} = 0011000\dots00_2$$

$$\text{EncodedExponent} = 2 + \text{Bias} = 2 + 127 = 129 = 10000001_2$$

0 10000001 001 1000 0000 0000 0000 0000

0x4098 0000

Single-Precision Range (Normal Numbers)

- In **normal** SP FP numbers, encoded exponents are in $[1, 254]$
 - 00000000_2 and 11111111_2 are reserved
- What is the smallest positive value of normal SP FP numbers?
- What is the largest positive value of normal SP FP numbers?

0 | 0 0000001 | 0000 ... 00

0 | 1 1111110 | 1 1111 ... 111

1	8	23
---	---	----

Single-Precision Range (Normal Numbers)

- In **normal** SP FP numbers, exponents are from 1 to 254
 - 00000000_2 and 11111111_2 are reserved

- Smallest positive value

- Exponent: $00000001_2 \Rightarrow \text{actual exponent} = 1 - 127 = -126$

- Fraction: $000\dots00 \Rightarrow \text{significand} = 1.0$

$$1.0 \times 2^{-126} \approx 1.2 \times 10^{-38}$$

How do we represent 0.0?

- Largest positive value

- Exponent: $11111110_2 \Rightarrow \text{actual exponent} = 254 - 127 = 127$

- Fraction: $111\dots11 \Rightarrow \text{significand} \approx 2.0$


$$2.0 \times 2^{+127} \approx 3.4 \times 10^{+38}$$

Denormalized/subnormal Numbers

- Denormalized number: the exponent field is 0
 - The actual exponent is always -126 for single precision numbers
 - The hidden bit is 0

$$v = (-1)^S \times (\text{0.Fraction}) \times 2^{-126}$$

- Denormalized numbers can represent numbers smaller than normal numbers
 - Allow for gradually approaching to 0, with diminishing precision



$$0.000000000000000000000001 \times 2^{-126} = 2^{-129}$$

Representation of 0

- 0 is a denormalized number !

All bits in exponent and fraction are 0.

But the sign can be 0 or 1. So we have two 0's!

0	0000 0000	000 0000 0000 0000 0000
---	-----------	-------------------------

1	0000 0000	000 0000 0000 0000 0000
---	-----------	-------------------------

$$x = (-1)^S \times (0.0) \times 2^{-126} = \pm 0.0$$

Infinites and NaN

Exponent = 1111 1111 (255)

- If fraction = 000...0
 - \pm Infinity
 - Can be used in subsequent calculations, avoiding need for overflow check
- If fraction \neq 000...0
 - Not-a-Number (NaN)
 - Indicates illegal or undefined result
 - e.g., 0.0 / 0.0
 - Can be used in subsequent calculations

Try these in Python:

```
float('inf') + 1.0  
float('inf') + float('-inf')
```

IEEE Floating-Point Format: double precision

single: 8 bits

double: 11 bits

single: 23 bits

double: 52 bits



Sign

0 : non-negative,
1: negative

Exponent in

excess-k representation

Bits after the binary point
There is a hidden 1 !

$$\text{value} = (-1)^S \times (1.\text{Fraction}) \times 2^{(\text{EncodedExponent} - \text{Bias})}$$

Exponent in single-precision: **excess-127**: Bias = 127.

Exponent in double-precision: **excess-1023**: Bias = 1023

Single precision vs double precision


IEEE FP

	Single	Double
Total number of bits	32	64
Number of bits in exponent	8	11
Number of bits in fraction	23	52
Bias	127	1023
Smallest positive value (normal values)	1.0×2^{-126} $\approx 1.18 \times 10^{-38}$	1.0×2^{-1022} $\approx 2.2 \times 10^{-308}$
Largest positive value	$2.0 \times 2^{+127}$ $\approx 3.4 \times 10^{+38}$	$2.0 \times 2^{+1023}$ $\approx 1.8 \times 10^{+308}$
Precision	23 bits ≈ 6 dec. digits	52 bits ≈ 16 dec. digits

F and D Extensions in RISC-V

RISC-V I32

- F for float and D for double
 - D is a superset. If D is supported, F is supported
- Separate FP register file (RF) consisting of 32 FP registers
 - ~~In F~~, each register can hold a float
 - In D, each register can hold a float or a double



f0, f1, ... f30, f31

f0 is not a special register

- FP instructions operate only on FP registers
 - Programs generally don't do integer ops on FP data, or vice versa
 - More registers with minimal code-size impact

FP register name and calling convention

FP Registers	Name	Usage
f0 - f7	ft0 - ft7	FP temporary registers. Not preserved
f8 - f9	fs0 - fs1	Callee saved registers. Preserved
f10 - f11	fa0 - fa1	First 2 arguments. Return values. Not preserved
f12 - f17	fa2 - fa7	6 more arguments. Not preserved
f18 - f27	fs2 - fs11	Callee saved registers. Preserved
f28 - f31	ft8 - ft11	FP temporary registers

12 callee saved registers. 12 temporary registers. 8 argument registers.

Load/store for FP numbers

- FP load and store instructions
 - w for SP and d for DP

flw, fsw, fld, fsd

same memory addressing modes

base address is an integer?

flw f8, 0(sp) # single-precision

fsw f8, 4(sp)

fld f9, 8(s1) # double-precision

fsd f9, 16(s1)

FP Arithmetic

- Single-precision arithmetic

fadd.s, fsub.s, fmul.s, fdiv.s, fsqrt.s

f0 = f1 + f6

fadd.s f0, f1, f6

Same as Int

- Double-precision arithmetic

fadd.d, fsub.d, fmul.d, fdiv.d, fsqrt.d

f1 ~~⋈~~ f2 * f3

fmul.d f1, f2, f3

32 32

32? X

Int Mult

Lower 2

higher 32

FP Comparison and Branch

- Single- and double-precision comparison

`f.eq.s, f.lt.s, f.le.s`

`f.eq.d, f.lt.d, c.le.d`

- Result, 0 or 1, is saved in **an integer destination register**
 - Use beq or bne to branch on comparison result

`# if f3 < f4, goto loop`

`f.lt.d t0, f3, f4`

`bne t0, x0, loop`

`# t0 = f3 < f4`

`# if t0`

Compare with x0
No need to compare with 1

Floating point precision

- Be mindful when you compare two FP numbers for equal

$$\underline{0.1 * 3} \neq \underline{0.3}$$

$$(0.1 * 3) - 0.3 = 5.55111512312578E-17$$

Handwritten notes: "≠ 0" above the result, "10⁻¹⁷" above the exponent, and a small "0" below the minus sign.

python

```
>>> float.hex(0.1)
'0x1.999999999999ap-4'
>>> float.hex(0.1*3)
'0x1.3333333333334p-2'
>>> float.hex(0.3)
'0x1.3333333333333p-2'
```

Handwritten notes: "a1" with an arrow pointing to the 'a' in the first hex string, and a blue underline under the second hex string.

Associativity

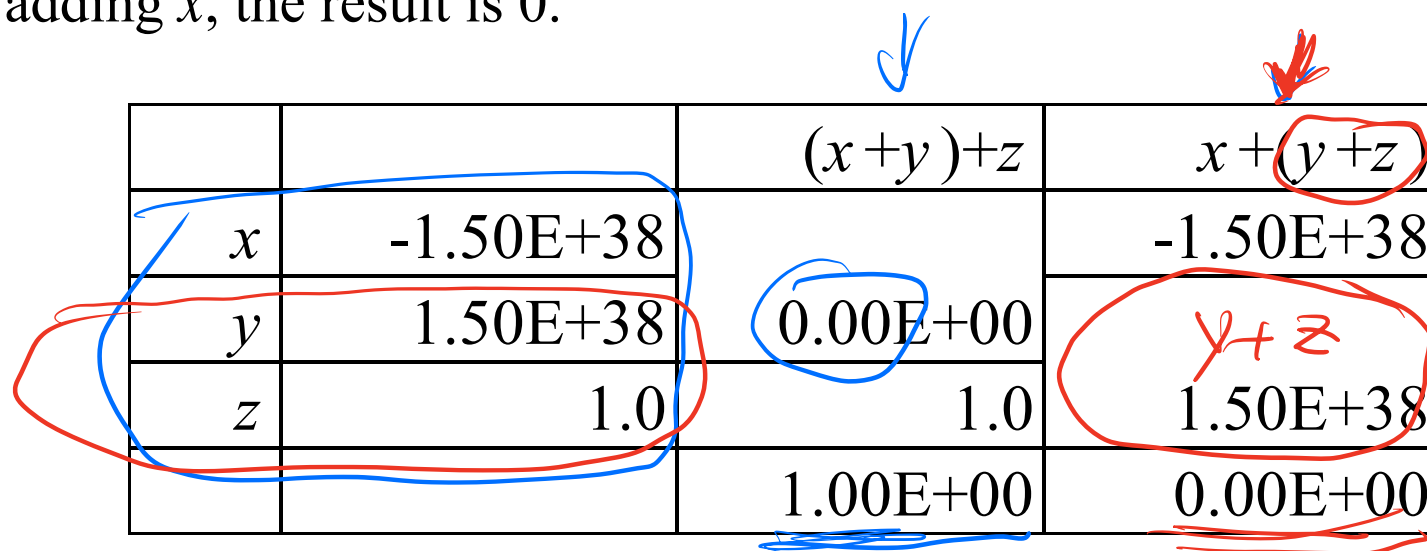
- Parallel programs may interleave operations in unexpected orders
 - Assumptions of associativity may fail
 - Need to validate parallel programs under varying degrees of parallelism

Example

$$(x + y) + z \neq x + (y + z)$$

If $(x + y)$ is computed first, the result is 0. After adding z , the result is 1.

If $(y + z)$ is computed first, the result is y (because y is much larger than z). After adding x , the result is 0.



		$(x + y) + z$	$x + (y + z)$
x	-1.50E+38		-1.50E+38
y	1.50E+38	0.00E+00	$y + z$ 1.50E+38
z	1.0	1.0	1.50E+38
		1.00E+00	0.00E+00

FP Example: °F to °C

C code:

```
float f2c (float fahr)
{
    return ((5.0/9.0)*(fahr - 32.0));
}
```

fahr in f10, return value in f10.

Constants 5.0, 9.0, and 32.0 are stored in (global) memory.

gp →

32.0
9.0
5.0

RISC-V code:

```
f2c: flw      f0, 0(gp)      # load 5
      flw      f1, 4(gp)      # load 9
      fdiv.s   f0, f0, f1     # compute 5/9
      flw      f1, 8(gp)      # load 32
      fsub.s   f10, f10, f1  # compute fahr - 32
      fmul.s   f10, f0, f10  # multiply with 5/9
      jalr     x0, 0(ra)
```


Frequency of RISC-V instructions in SPEC CPU2006

Figure 3.22

17 most popular instructions
76% of all instr. executed

RISC-V Instruction	Name	Frequency	Cumulative
Add immediate	<u>addi</u>	14.36%	14.36%
Load doubleword	ld	8.27%	22.63%
Load fl. pt. double	fld	6.83%	29.46%
Add registers	<u>add</u>	6.23%	35.69%
Load word	<u>lw</u>	4.38%	40.07%
Store doubleword	sd	4.29%	44.36%
Branch if not equal	<u>bne</u>	4.14%	48.50%
Shift left immediate	<u>slli</u>	3.65%	52.15%
Fused mul-add double	fmadd.d	3.49%	55.64%
Branch if equal	<u>beq</u>	3.27%	58.91%
Add immediate word	addiw	2.86%	61.77%
Store fl. pt. double	fsw	2.24%	64.00%
Multiply fl. pt. double	fmul.d	2.02%	66.02%
Load upper immediate	lui	1.56%	67.59%
Store word	<u>sw</u>	1.52%	69.10%
Jump and link	<u>jal</u>	1.38%	70.49%
Branch if less than	<u>blt</u>	1.37%	71.86%
Add word	addw	1.34%	73.19%
Subtract fl. pt. double	fsub.d	1.28%	74.47%
Branch if greater/equal	<u>bge</u>	1.27%	75.75%

Summary

- Support for data types and arithmetic are part of ISA design
- RISC-V
 - Base supports integer add and sub
 - M extension supports mul and div
 - F and D extensions support FP operations
- Exceptions during arithmetic
 - Operations can overflow
 - Need to handle error with hardware and/or software
 - Floating-point has bounded range and precision
- Bits can be interpreted in many ways
 - Signed, unsigned, instruction, characters, FP numbers

Denormalized Numbers Examples

In the table, only the first number is a normal number

Exponent	Fraction	Actual exponent in decimal	Value
0000 0001	00000...00	-126	1.0×2^{-126} (normal number)
0000 0000	10000...00	-126	$0.1 \times 2^{-126} = 2^{-127}$
0000 0000	01000...00	-126	$0.01 \times 2^{-126} = 2^{-128}$
...			
0000 0000	00000...01	-126	$0.0...01 \times 2^{-126} = 2^{-149}$
0000 0000	00000...00	-126	$0.0...00 \times 2^{-126} = 0$

Conversion between datatypes

- Many conversion instructions. Study the reference card

`fcvt.s.w, fcvt.d.w, fcvt.d.s, ...`

```
addi      t0, x0, 5
fcvt.s.w   ft0, t0      # word to single-precision
fcvt.d.w   ft1, t0      # word to double-precision
# ft0 is a single-precision 5.0
# ft1 is a double-precision 5.0
```

Loading constants from memory:

[cse3666/91-f2c.s at master · zhijieshi/cse3666 \(github.com\)](#)

Using conversion instructions:

[cse3666/91-f2c-v2.s at master · zhijieshi/cse3666 \(github.com\)](#)

Question

Convert the decimal number 0.9 to a binary number

0						
2^0	2^{-1}	2^{-2}	2^{-3}	2^{-4}	2^{-5}	2^{-6}

Converting decimal to binary Example

Decimal	Binary
0.9	0.
$0.9 * 2 = 1.8$	0.1
$0.8 * 2 = 1.6$	0.11
$0.6 * 2 = 1.2$	0.111
$0.2 * 2 = 0.4$	0.1110

We can find the first 4 digits after the binary point by the following steps:

$$0.9 * 2^4 = 14.4$$

Convert 14 to 4-bit binary number and we get 1110.

Example: Convert to Single-Precision FP numbers

Represent -0.75 with a single precision floating-point number

$$-0.75 = -0.11_2 = (-1)^1 \times 1.1_2 \times 2^{-1}$$

$$S = 1$$

$$\text{Fraction} = 1000\dots00_2$$

$$\text{EncodedExponent} = -1 + \text{Bias} = -1 + 127 = 126 = 01111110_2$$

1 01111110 100 0000 0000 0000 0000 0000

0xBF40 0000

Reading Single-Precision FP Number - Solutions

0x C1C0 0000

1100 0001 1100 0000 0000 0000 0000 0000

$S = 1$

Fraction = $10000\dots00_2$

Encoded Exponent = $10000011_2 = 131$ (as unsigned)

Actual exponent = $131 - 127 = 4$

The value is

$$\begin{aligned} & (-1)^1 \times (1 + 0.1_2) \times 2^{(131 - 127)} \\ &= -1 \times 1.5 \times 2^4 \\ &= -24 \end{aligned}$$